



Bits on Bits: Showcasing Next-Gen Arithmetic through Information Theory

Max Hawkins and Rich Vuduc

Computing Export Controls

“Who controls the ~~spice~~ compute, controls the world”



Mar 24, 2025 | Hudson Institute

AI, National Security, and the Global Technology Race: How US Export Controls Define the Future of Innovation



Nury Turkel

Exclusive: Nvidia modifies H20 chip for China to overcome US export controls, sources say

By Liam Mo and Brenda Goh

May 9, 2025 5:21 AM EDT · Updated May 9, 2025

Definitions of computing performance impact billion-dollar decisions, national security, and the future of computing.

How to do this accurately, fairly, and generally?

Which computer is more performant?

By how much?

Computer A

- 1.44 Exaflop/s*
- Nvidia's GB200 NVL72
 - "The NVIDIA GB200 NVL72 is an exascale computer in a single rack." -



Computer B

- 1.35 Exaflop/s**
- Frontier
 - #2 most performant public supercomputer




* With sparse FP4 tensor cores

** Dense FP64 with Linpack

The Variety of Arithmetic

- Data types
 - Bit widths
 - Bit allocation (e.g. mantissa and exponent bits)
 - Encoding schemes - integer vs floats vs posits...
 - Specifications (e.g. IEEE, OCP, vendor...)
- Operations
 - Add, subtract, negate, multiply, divide, compare, sqrt, tanh, ...
 - Sparsity
 - Emulation (Ozaki and beyond)
 - Noise
 - Scalar vs vector vs matrix inputs



Source: Nvidia

Hardware:
Quantum, analog,
neuromorphic, reversible, ...

How do we fairly measure and compare performance
across this large design space?

+ .
o

What are we doing now?

What could we do now?

What can we do in the future?

+ .
o

What are we
doing now?

+

•

◦


+

•

◦

Weighting Data Types by Bit Widths

- FP4 vs FP64
 - ‘Exaflop’ computers: Frontier vs GB200 NVL72
 - 64 bits $\rightarrow 2^{64}$ possible states
 - 4 bits $\rightarrow 2^4$ possible states
- Linear comparisons?
 - $\frac{2^{64}}{2^4} = 1,152,921,504,606,846,976$
- Logarithmic comparisons?
 - $\frac{\log_2(2^{64})}{\log_2(2^4)} = \frac{64}{4} = 16$
- We use logarithms of the state space to compare across bit widths
 - Bit width approximation
- U.S. Gov’t export controls use this approach



Linear
Weighting




Logarithmic
Weighting

imgflip.com


Reducing Redundant Encodings

- How many **distinct** states does a data type represent?
- Redundancy wastes bits/bitstrings
 - “There should be no redundant bit patterns to mean the same thing; every bit counts.” – John Gustafson




Optimizing Data Type Usage Efficiency

- “Does this kernel need FP64, or can I use FP16, Int8, FP2,...?”
 - Much existing bit inefficiency!
- Value range
 - Most data spans << 300 decades
- Accuracy
 - Many applications don’t always need 53 bits of relative accuracy
 - Even HPL (see HPL-MxP)



Innovating in Data Types and Emulation

- Block-scaled encodings
- Posits, takums, ...
- Ozaki emulation
 - Performing floating-point matmul with lower-precision hardware
 - Useful when:
 - High-precision performance is low
 - Data spans a very small range and requires little accuracy




What are we doing now?

- Weighing data types by their bit widths
 - Log scaling of state space
- Reducing redundant value encodings
- Optimizing data type usage efficiency with smaller data types
 - “Does this kernel really need FP64?”
- Innovating in data types and emulation
 - Block-scaled FP, Posits, Takums, Ozaki emulation, and beyond

What could we
do now?

Shannon Entropy in Brief

- Uncertainty
- Flipping a coin: Heads or Tails \rightarrow 1 bit
- Rolling a die with M faces $\rightarrow \log_2(M)$
- Shannon entropy (H): A measure of uncertainty
 - Discrete random variable X with probability distribution $p(x)$
 - Measured in bits if $b = 2$



$$H(X) = - \sum_{x \in \mathcal{X}} p(x) \log_b p(x)$$

Reframe Effects of Redundant Encodings with Entropy

- Redundancy reduces *information capturing potential*
 - Magnified by smaller number of bitstrings – low bit widths
- Quantified with encoding efficiency: $\eta = \frac{H_{encoding}}{bit_width}$
 - Careful when mixing linear and log scaling
 - 50% bitstring redundant 4-bit encoding → 3 bits of entropy (not 2!)
- Example: IEEE-754 redundant NaN encodings

Every physical **informational** bit counts!

Encoding Efficiency in Practice (NaNs only)

Float Type	Bit Width	Encoding Entropy	Encoding Efficiency (%)
IEEE-754	64	63.974	99.960
IEEE-754	32	31.906	99.707
TF32	19	18.957	99.774
IEEE-754	16	15.657	97.854
BF16	16	15.969	99.806
OCP (E4M3)	8	7.992	99.902
IEEE-754 (E4M3)*	8	7.792	97.397
OCP (E2M3)	6	6	100
IEEE-754 (E2M3)*	6	5.167	86.119
IEEE-P3109	Any	Ideal	100
Posits	Any	Ideal	100
Takums	Any	Ideal	100

* Theoretical – Does not exist.

Data Type Usage Efficiency and Info Theory

- Bits are your currency, and you allocate them as needed
 - “Do I really need FP64 for everything?”
- How to answer that analytically: Info theory
 - Bitstrings for values beyond needed range go unused → 0 entropy
 - Bits allocated towards excess precision are baggage
 - Constant values have 0 uncertainty → 0 entropy




- Example: Using FP8 (E4M3) to encode 4-bit integer
 - Bit width approximation: 8 physical bits per op to 4 → 2x performance ‘reduction’
 - Info Theory: 4 bits of info to 4 → No performance impact

Info theory doesn’t ‘punish’ optimization (unlike bit width approx.)

Another Problem with Mixed Bit Field Mentality

- Float encodings unnecessarily separate value range and precision!
 - Ex: BF16 vs FP16 or FP8 E4M3 vs E3M4
- In a purely log or purely integer format, these differences disappear
 - Can exactly tradeoff accuracy \leftrightarrow range
 - Simplifies reasoning about data types





■ Log 3bit (1-8) ● Log 4bit (1-16) ▲ Log 4bit (1-8) ♦ Log 4bit (1-16 rescaled)



From Quantization/Communication to Operations

How much does a given operation *reduce uncertainty*?

- Constant inputs or outputs → **None!**
 - Compile them away
- More output states → **More!**
- Some states more likely → **Less!**
 - NaNs, overflow, underflow, etc
- Noisy/Error-prone HW → **Less!**
 - Determinism matters



From Quantization/Communication to Operations

"The fundamental problem of communication is that of reproducing at one point either exactly or approximately a message selected at another point"

– *A Mathematical Theory of Communication (Shannon 1948)*

- Generalized the concept of communication performance
 - Allowed for fair and generalized performance evaluation







Fig. 1—Schematic diagram of a general communication system.

Communication: Copies inputs to outputs (identity operation)




Expand beyond the identity: Computation



Mutual Information (I)

- Measures the ‘shared’ information of random variables
 - Considers properties of the operation and noise
- ‘Operational’ quantity
- Application and runtime-dependent

$$I(X; Y) = H(X) - H(X|Y)$$



Channel Capacity (C)

- Maximum MI over possible distributions
- Upper bound/ideal quantity
- Application-agnostic

$$C = \max_{p(x)} I(X; Y)$$


Info Theory and Computation

- Expand Shannon's communication performance model
 - Identity operator (communication) → Arbitrary operator (computation)
 - Mutual info and channel capacity naturally handle this
- Flop/s → Bit/s
 - Base Measure: Uncertainty reduction
 - Operational: Mutual Information
 - Ideal/Peak: Channel capacity
- Enable generalized and fair performance evaluation
 - Just like communication has had for >70 years
- Aligns with existing performance metrics

Every *informational* bit counts
(for communication, quantization, and computation)

Data Type Channel Capacity Estimation


- Estimate the channel capacity of operations
- Inputs: Every possible combination of 8-bit values
- Outputs: Add/Mult in 8 or 16-bit formats



Incorporate Data Type Utilization Into Tools

Information Roofline

- Extends traditional roofline
- Adds data type utilization
 - Needed with today's innovation
 - Every bit counts!
- X-Axis: Arithmetic Intensity
 - Unitless $\left(\frac{bits_{comp}}{bits_{comm}}\right)$
- Y-Axis: Information Throughput
 - Bits of information per second




Either you're inefficiently using a data type, or you traded accuracy/range for op-throughput.
The information roofline makes this explicit.

Entropy Staircase

- Simplify complex execution choices
 - Native pipelines and emulation
- Quantify benefits of emulation
- Ease of ascending/descending the entropy staircase
 - Posits vs floats
 - Traditional ‘mixed’ precision

GEMM on Nvidia H200 GPU



What could we do now?

- Reframe redundant value encodings with entropy loss
- Measure data type and operation performance with info. theory
 - Unifies communication and computation
- Incorporate data type utilization into usable tools
 - Information roofline
 - Entropy staircase

What could we
do in the future?

Uncorrectable Noise



This is real!

H100 GPU in space.

- Status quo: Digital computing is nearly lossless
 - Data type designers: “We don’t need to care about bit flips.”
 - Infs/NaNs/bitfields make bitflips disastrous
- Motivations for accepting error:
 - Undervolting: Energy consumption \leftrightarrow Error tradeoff
 - Space-based datacenters: Radiation causing bit flips
- Shannon capacity of graphs
 - Single bit flips: Hypercube
- Could digital encodings look more like neuromorphic spike trains?

[Image Source](#)

What could we do next?

- Noise/error-tolerant formats and algorithms
- Variable-width encodings
- Dataflow optimization (hyper-localized data types/operations)
- Compare performance across hardware paradigms
 - Quantum
 - Neuromorphic
 - Analog
 - Reversible

There is still room for innovation.

When will the juice be worth the squeeze?

How will we fairly and generally measure computing performance in the future?

- Innovation will continue
 - Loss of Moore's law/Dennard scaling
- Data type and hardware variety will grow
- How many asterisks is too many?
 - *Flop/s in FP32, with sparsity, using tensor cores, emulated with Ozaki scheme 1 using two slices of Int8
- Export controls need a fundamental grounding
- HPC + Quantum/Neuromorphic/Analog/Reversible systems

We need to innovate in performance measurement and tooling to capture the evolving hardware and data types.

Information Theory
Enables a Useful and
General Framework
for Computing
Performance

Every *informational* bit counts!

4-page preprint:
[arxiv.org/pdf/2508.05621](https://arxiv.org/pdf/2508.05621.pdf)

mhawkins60@gatech.edu

richie@cc.gatech.edu



Bonus Slides!

+

•

○

+

•


○


Aspects Information Theory Explicitly Ignores

- Ease of physical implementation
- Semantics and error analysis
- Reproducibility

Max Relative Error and Logarithmic Encodings

- Logarithmic encodings are optimal*
 - ...if using max relative error on a bounded interval





(Computer) Arithmetic

- What is the ‘freedom’ or set of potential values of variable x ?
 - $x \in R$
- Mathematician: If $x \in R$, infinitely many values!
- Hardware designer: If stored in FP64, $\sim 2^{64}$ states
- HPC practitioner: $[-1000, 1000]$ but mostly close to zero
- “All models are wrong. Some are useful”